

Chapter VI


Vertex Shader

GPU Rendering Pipeline (revisited)

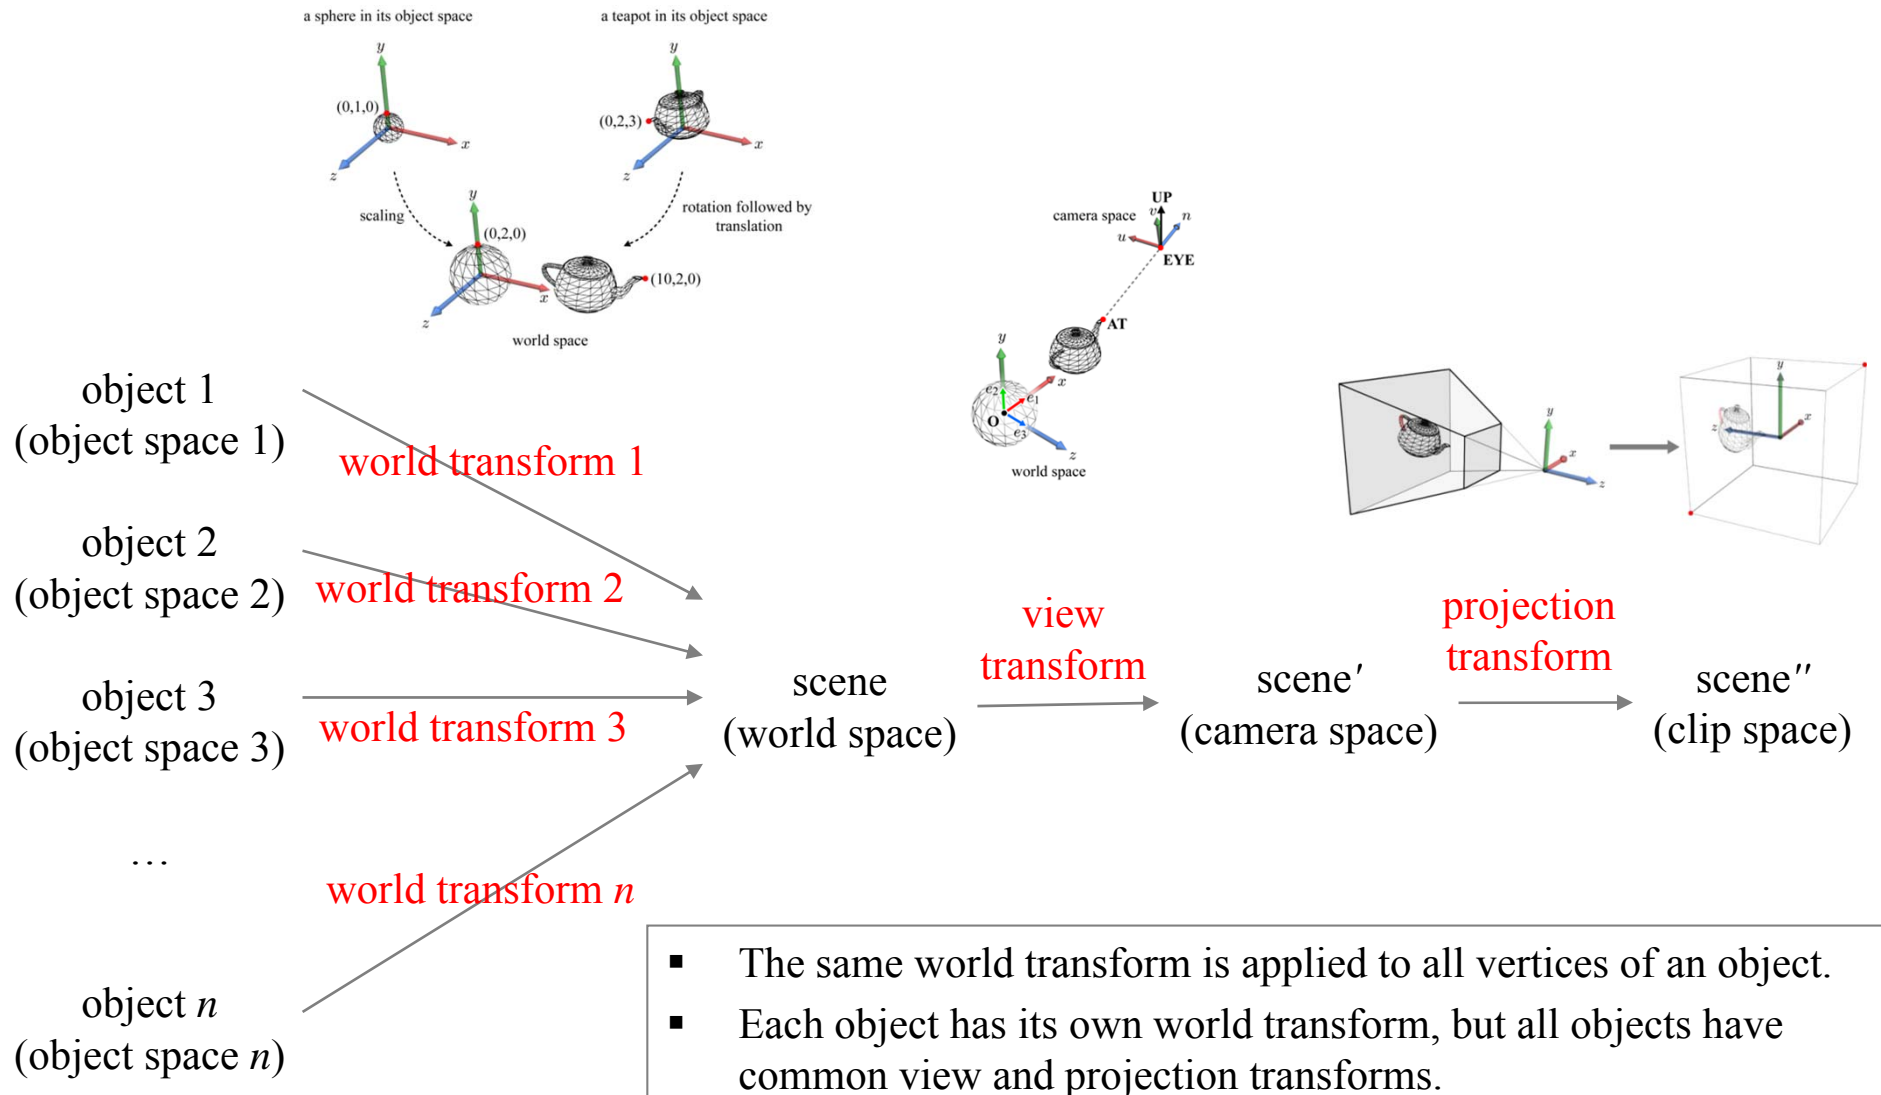
- Main stages in the *rendering pipeline*



- The vertex shader (vertex program) operates on every input vertex stored in the vertex array and performs various operations.
- The essential among them is applying a series of *transforms* to the vertices.




Transforms (revisited)



Vertex and Index Arrays (revisited)

- The indispensable components of vertex array are position, normal, and texture coordinates. Texture coordinates are 2D and used to access the 2D texture. (This is presented at CH 8.)
- Conceptually, the vertex shader processes a vertex (an element) of the vertex array at a time.



vertex array			index array	
0	(0.000, 1.000, 0.000)	(0.000, 1.000, 0.000)		0
1	(0.000, 0.707, 0.707)	(0.000, 0.663, 0.748)		1
2	(0.500, 0.707, 0.500)	(0.529, 0.663, 0.529)		2
3	(0.000, 0.000, 1.000)	(0.000, 0.000, 1.000)		1
4	(0.707, 0.000, 0.707)	(0.707, 0.000, 0.707)		3
	.	.		4
	.	.		4
	.	.		2
	.	.		1
	.	.		.
	.	.		.
25	(0.000, -1.000, 0.000)	(0.000, -0.707, -0.707)		143
	position	normal	tex coord	16

OpenGL ES

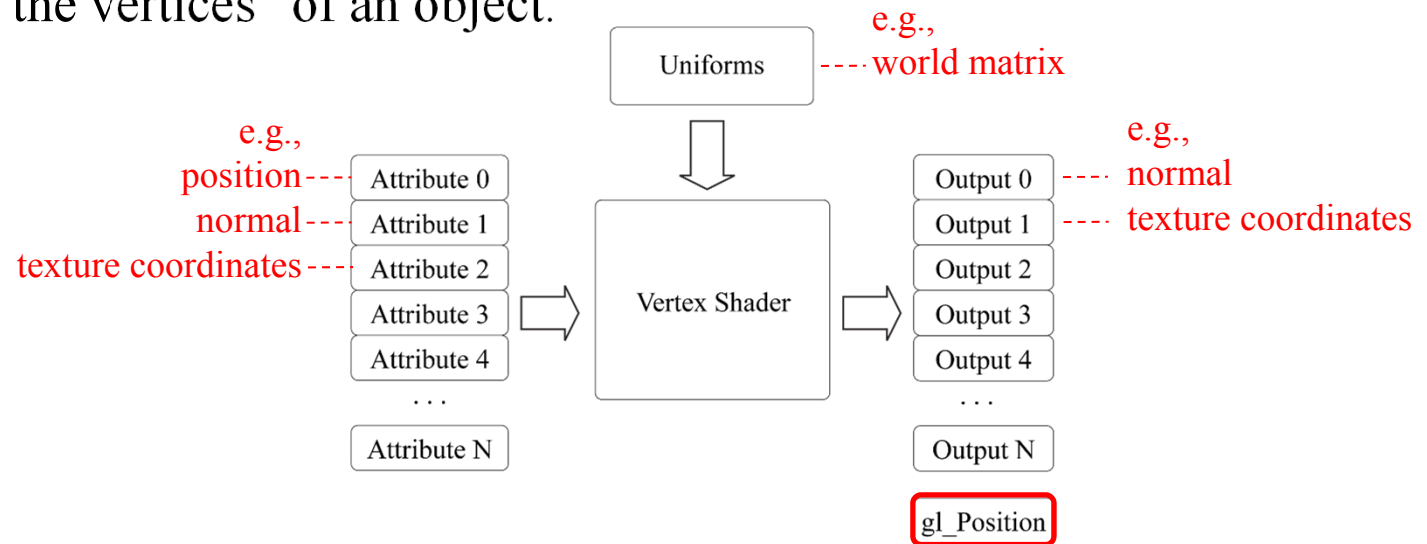
- OpenGL ES is a subset of OpenGL.
 - OpenGL ES 2.0 is derived from OpenGL 2.0 and provides vertex and fragment shaders.
 - OpenGL ES 3.0 is derived from OpenGL 3.3 and adds many enhanced features to OpenGL ES 2.0.
 - OpenGL ES 3.1 (publicly released in March 2014) includes compute shaders.
 - OpenGL ES 3.2 (August 2015) includes geometry and tessellation shaders.
- OpenGL ES 3.0 is backward compatible with OpenGL ES 2.0, i.e., applications built upon OpenGL ES 2.0 work with OpenGL ES 3.0.
- This class focuses on OpenGL ES 3.0.
 - OpenGL ES 3.0 API specification
 - OpenGL ES Shading Language Specification
- From now on, let's call OpenGL ES simply 'GL' and OpenGL ES Shading Language by 'SL.'

OpenGL ES Shading Language

- SL is a C-like language. For example, `float` is supported.
- However, SL works on GPU, the goal and architecture of which are different from those of CPU. The differences shape SL in a distinct way.
- Vector and matrix examples
 - `vec4` defines a floating-point 4D vector and `ivec3` defines an integer 3D vector.
 - `mat3` and `mat4` define 3x3 and 4x4 ‘square’ matrices, respectively, whereas `mat3x4` is for a 3x4 matrix. The matrix elements are all `float` values.

Vertex Shader

- A vertex shader per object!
- The main input
 - Attributes: *Per-vertex* data that are typically provided in *vertex array*.
 - Uniforms: Read-only values such as the transform matrix to be uniformly applied to the vertices “of an object.”



- The output
 - They must include the built-in variable, `gl_Position` that stores the clip-space vertex position.
 - In addition, they usually include normal, texture coordinates, etc.

Vertex Shader (cont'd)

- Our first vertex shader

```
#version 300 es

uniform mat4 worldMat, viewMat, projMat;

layout(location = 0) in vec3 position;
layout(location = 1) in vec3 normal;
layout(location = 2) in vec2 texCoord;

out vec3 v_normal;
out vec2 v_texCoord;

void main() {
    gl_Position = projMat * viewMat * worldMat * vec4(position, 1.0);
    v_normal = normal;
    v_texCoord = texCoord;
}
```

- Attributes are read-only values, which require the keyword, `in`.
 - The layout qualifiers such as `layout(location = 0)` specify the *attribute locations*. (The layout qualifiers are optional though.)
-

GL Program

- The fragment shader is written in a similar fashion.
- While the vertex and fragment shaders are in charge of low-level details in rendering, the role of GL program itself is managing the shaders and various data needed for the shaders.
- GL API
 - GL commands begin with the prefix `gl`.
 - GL data types begin with the prefix `GL`.

GL Program - Shader Object

- Given a vertex shader stored in a file, we do the following:
 - Its source code is loaded. Assume that it is pointed by `source` that is of type `char*`.
 - A new *shader object* is created using `glCreateShader`, which takes either `GL_VERTEX_SHADER` or `GL_FRAGMENT_SHADER` and returns the unsigned integer ID of the shader object.
 - Taking the vertex shader's source code and the shader object, `glShaderSource` stores the source code in the shader object.
 - The shader object is compiled using `glCompileShader`.

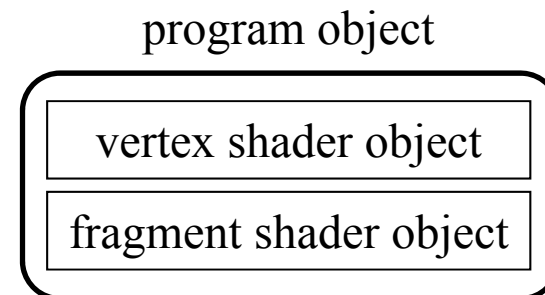
```
GLuint shader = glCreateShader(GL_VERTEX_SHADER);  
glShaderSource(shader, 1, &source, NULL);  
glCompileShader(shader);
```

- The same process has to be done for the fragment shader using `GL_FRAGMENT_SHADER` for `glCreateShader`.
 - Now you have two shader objects (for the vertex and fragment shaders) .
-

GL Program - Program Object

- The shader objects (for the vertex and fragment shaders) should be attached to a *program object*, which is then linked into the final executable.
 - The *program object* is created by `glCreateProgram`, which takes no argument and simply returns the ID of the new program object.
 - The shader and program objects are given to `glAttachShader`, which attaches the shader object to the program object.
 - Then, the program object is linked by `glLinkProgram`.
 - Finally, the program object will be used for rendering! For this purpose, `glUseProgram` is invoked.

```
GLuint program = glCreateProgram();  
glAttachShader(program, shader);  
glLinkProgram(program);  
glUseProgram(program);
```



- The above code is for vertex shader and so `glAttachShader` should be called with the fragment shader object.
-

Attributes

- The GL program not only hands over the attributes and uniforms to the vertex shader but also informs the vertex shader of their structures.
- Like the shader source codes, the mesh data stored in a file (such as .obj file) will be read into the vertex and index arrays of the GL program.
- Suppose that `vertices` and `indices` are the pointers to the arrays.

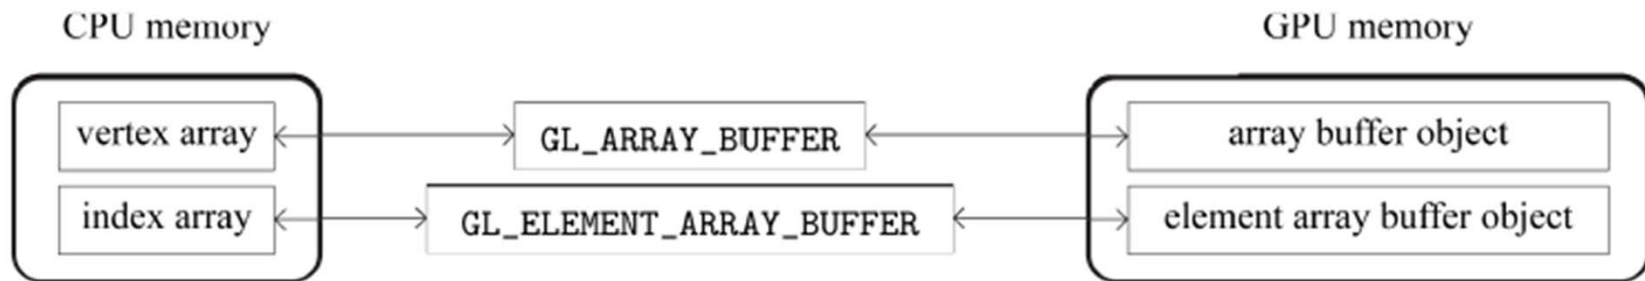
```
struct Vertex
{
    glm::vec3 pos; //position
    glm::vec3 nor; //normal
    glm::vec2 tex; //texture coordinates
};
```

```
std::vector<Vertex> vertices;
std::vector<GLushort> indices;
```

- `glm` stands for OpenGL Mathematics. It is a library that provides classes and functions with the same naming conventions and functionalities as SL.

Attributes (cont'd)

- The address space where the arrays reside is called the *client memory* by convention. The arrays in the client memory will then be transferred into what are called *buffer objects* in the GPU memory.
- For the indexed representation of a mesh, GL supports two types of buffer objects:
 - Array buffer object is for the vertex array and is specified by `GL_ARRAY_BUFFER`.
 - Element array buffer object is for the index array and is specified by `GL_ELEMENT_ARRAY_BUFFER`.



Attributes (cont'd)

- Creating and binding VBO
 - `glGenBuffers(GLsizei n, GLuint *buffers)` is asked for `n` buffer objects and returns them in `buffers`.
 - `glBindBuffer(GLenum target, GLuint buffer)` - `target` is either `GL_ARRAY_BUFFER` or `GL_ELEMENT_ARRAY_BUFFER`.
 - The buffer object is filled with data using `glBufferData(GLenum target, GLsizeiptr size, const void *data, GLenum usage)`
 - `target` is either `GL_ARRAY_BUFFER` or `GL_ELEMENT_ARRAY_BUFFER`.
 - `data` points to the vertex or index array supplied by the application. (This will be either vertices or indices.)

```
GLuint vbo;  
glGenBuffers(1, &vbo);  
glBindBuffer(GL_ARRAY_BUFFER, vbo);  
glBufferData(GL_ARRAY_BUFFER, vertices.size() * sizeof(Vertex),  
vertices.data(), GL_STATIC_DRAW);
```

Attributes (cont'd)


- The same process is done for the index array.

```
GLuint ibo;  
glGenBuffers(1, &ibo);  
glBindBuffer(GL_ELEMENT_ARRAY_BUFFER, ibo);  
glBufferData(GL_ELEMENT_ARRAY_BUFFER, indices.size() * sizeof(GLushort),  
indices.data(), GL_STATIC_DRAW);
```

- OK, now we have the VBO for vertex and index arrays.

Attributes (cont'd)

- The vertex array can be described as an *array of structures*, where a structure contains all attributes (pos, nor, and tex) of a vertex.



- The GL program uses `glEnableVertexAttribArray` and `glVertexAttribPointer` to inform the vertex shader of such a structure.

```
int stride = sizeof(Vertex);  
int offset = 0;
```

```
glEnableVertexAttribArray(0); // position = attribute 0  
glVertexAttribPointer(0, 3, GL_FLOAT, GL_FALSE, stride, (GLvoid*)offset);
```

```
offset += sizeof(glm::vec3); // for accessing normal  
glEnableVertexAttribArray(1); // normal = attribute 1  
glVertexAttribPointer(1, 3, GL_FLOAT, GL_FALSE, stride, (GLvoid*)offset);
```

```
offset += sizeof(glm::vec3); // for accessing texture coordinates  
glEnableVertexAttribArray(2); // texture coordinates = attribute 2  
glVertexAttribPointer(2, 2, GL_FLOAT, GL_FALSE, stride, (GLvoid*)offset);
```


Attributes (cont'd)

```
void          glVertexAttribPointer(GLuint index,
                                     GLint size,
                                     GLenum type,
                                     GLboolean normalized,
                                     GLsizei stride,
                                     const GLvoid * pointer);
```

index	the index of the vertex attribute
size	the number of components of the attribute, which must be 1, 2, 3, or 4
type	the data type such as <code>GL_INT</code> and <code>GL_FLOAT</code>
normalized	If <code>GL_TRUE</code> , integer data are mapped to <code>[-1,1]</code> or <code>[0,1]</code> before being used in the vertex shader.
stride	the byte distance between the consecutive attributes of the same kind
pointer	offset (in bytes) into the first occurrence of the attribute in the buffer object

Uniforms

- Our vertex shader has three uniforms: `worldMat`, `viewMat`, and `projMat`.
- Consider a dynamic environment.
 - If the scene object moves, `worldMat` should be changed.
 - If the viewpoint moves, `viewMat` should be changed.
- The GL program updates and provides them for the vertex shader.
- For this purpose, we first have to use `glGetUniformLocation` to find the uniform's 'location' in the program object, which has been determined during the link phase.
- In order to load the uniform, `worldMat`, for example, we use `glUniformMatrix4fv`, where 4 indicates a 4x4 matrix, `f` indicates that its elements are floating-point values, and `v` implies that the values are passed in a vector, i.e., an array:

```
glm::mat4 worldMatrix;  
GLint loc = glGetUniformLocation(program, "worldMat");  
glUniformMatrix4fv(loc, 1, GL_FALSE, &worldMatrix);
```

Uniforms (cont'd)

```
GLint  glGetUniformLocation(GLuint program,  
                             const GLchar *name);
```

program program object
name uniform name in a string

```
void    glUniformMatrix4fv(GLint location,  
                            GLsizei count,  
                            GLboolean transpose,  
                            const GLfloat *value);
```

location uniform location
count the number of matrices to be modified
transpose This must be GL_FALSE.
value pointer to an array of 16 GLfloat values

Drawcalls

- We have made all attributes and uniforms available.
- Suppose that we have a good fragment shader, whatever it is.
- Then, we can draw a polygon mesh.
- For rendering a polygon mesh, we can make a *drawcall*. We have two choices:
 - `glDrawArrays` for non-indexed mesh representation
 - `glDrawElements` for indexed mesh representation

```
void glDrawArrays(GLenum mode,  
                  GLint first,  
                  GLsizei count);
```

`mode` `GL_TRIANGLES` for polygon mesh
`first` start index in the vertex array
`count` the number of vertices to draw

Drawcalls (cont'd)

- For indexed representation:

```
void    glDrawElements(GLenum mode,  
                        GLsizei count,  
                        GLenum type,  
                        const GLvoid * indices);
```

mode GL_TRIANGLES for polygon mesh
count the number of indices to draw
type index type
indices byte offset into the buffer bound to GL_ELEMENT_ARRAY_BUFFER

```
glDrawElements(GL_TRIANGLES, indices.size(), GL_UNSIGNED_SHORT, 0);
```

- In the example, the index array has 144 elements and the vertex array has 26 elements. Therefore `indices.size()` returns 144.

